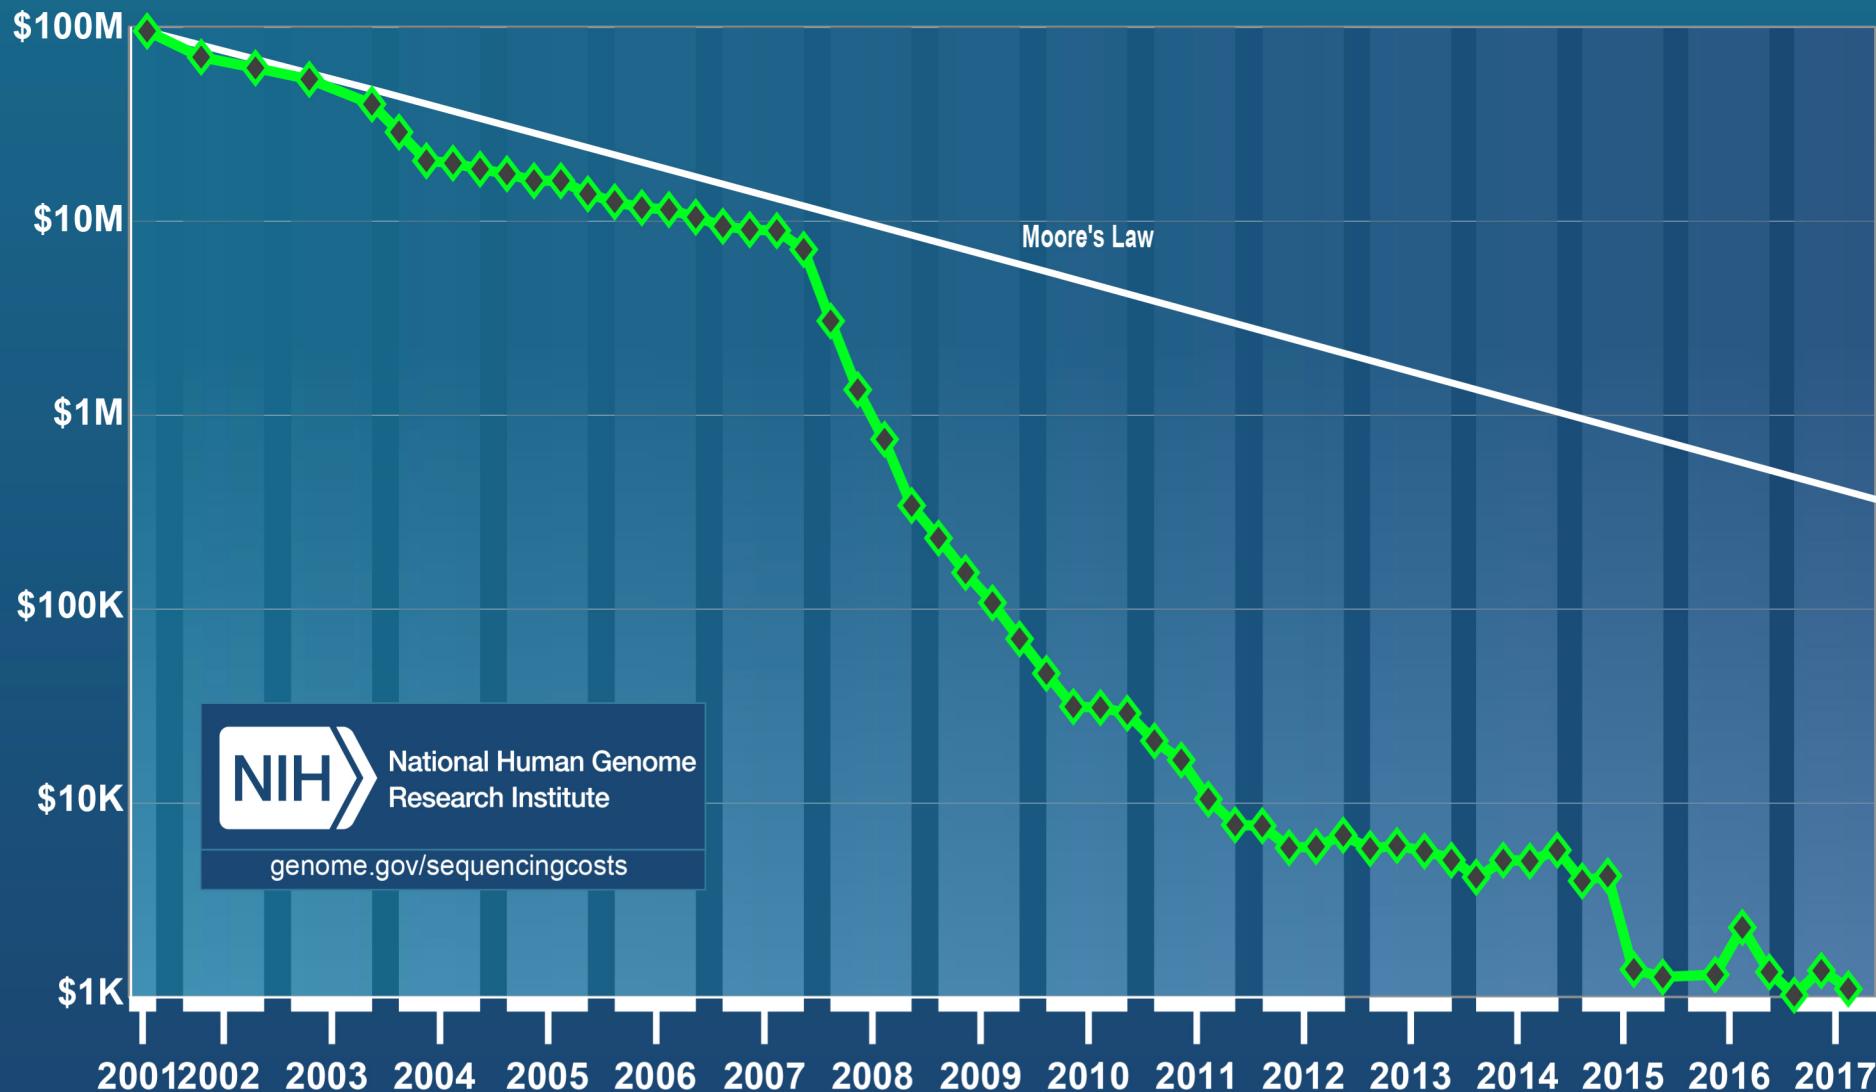


Genome Sequencing

Marc Tollis

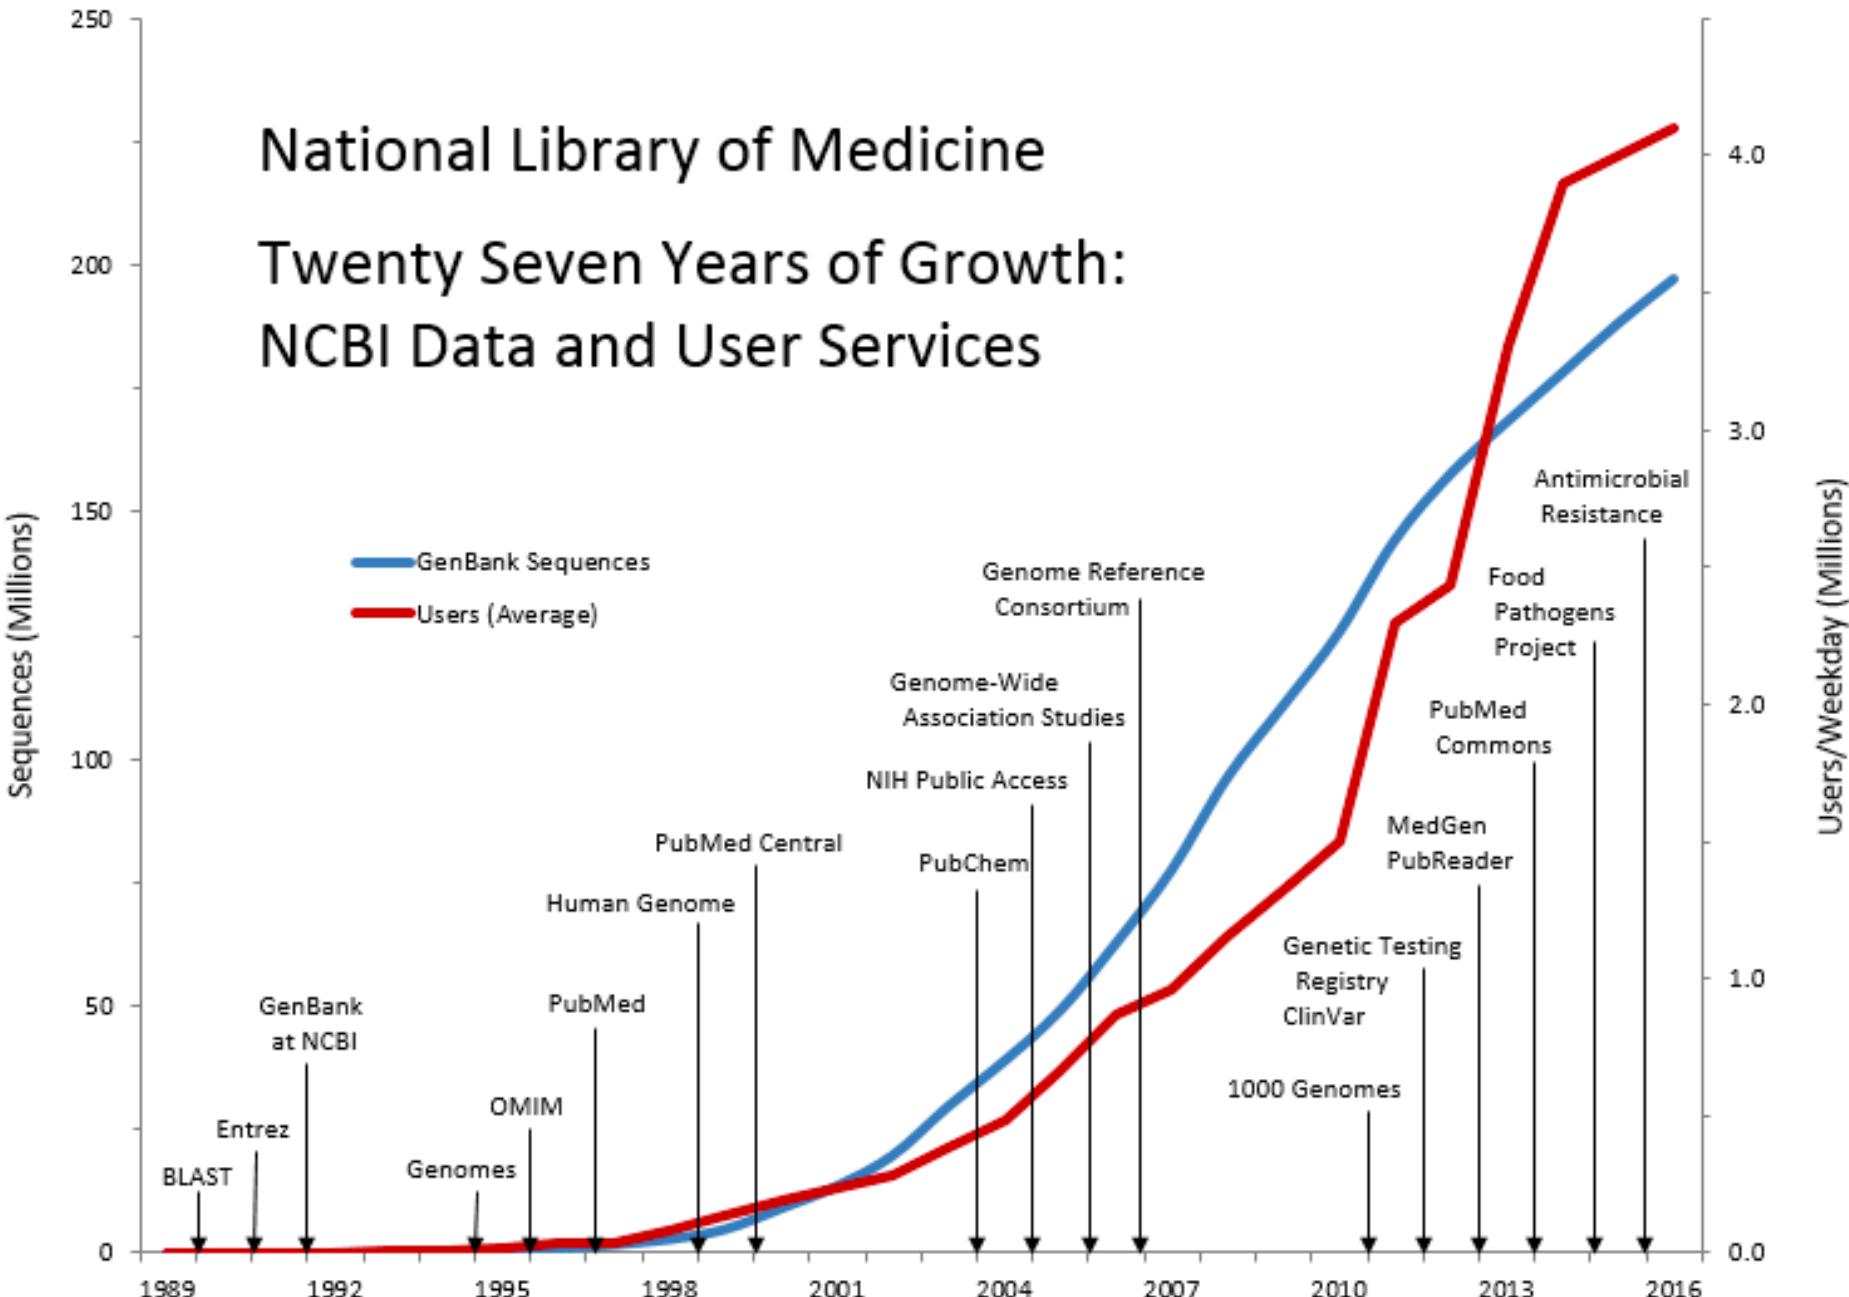
Comparative Genomics

Cost per Genome



National Library of Medicine

Twenty Seven Years of Growth: NCBI Data and User Services

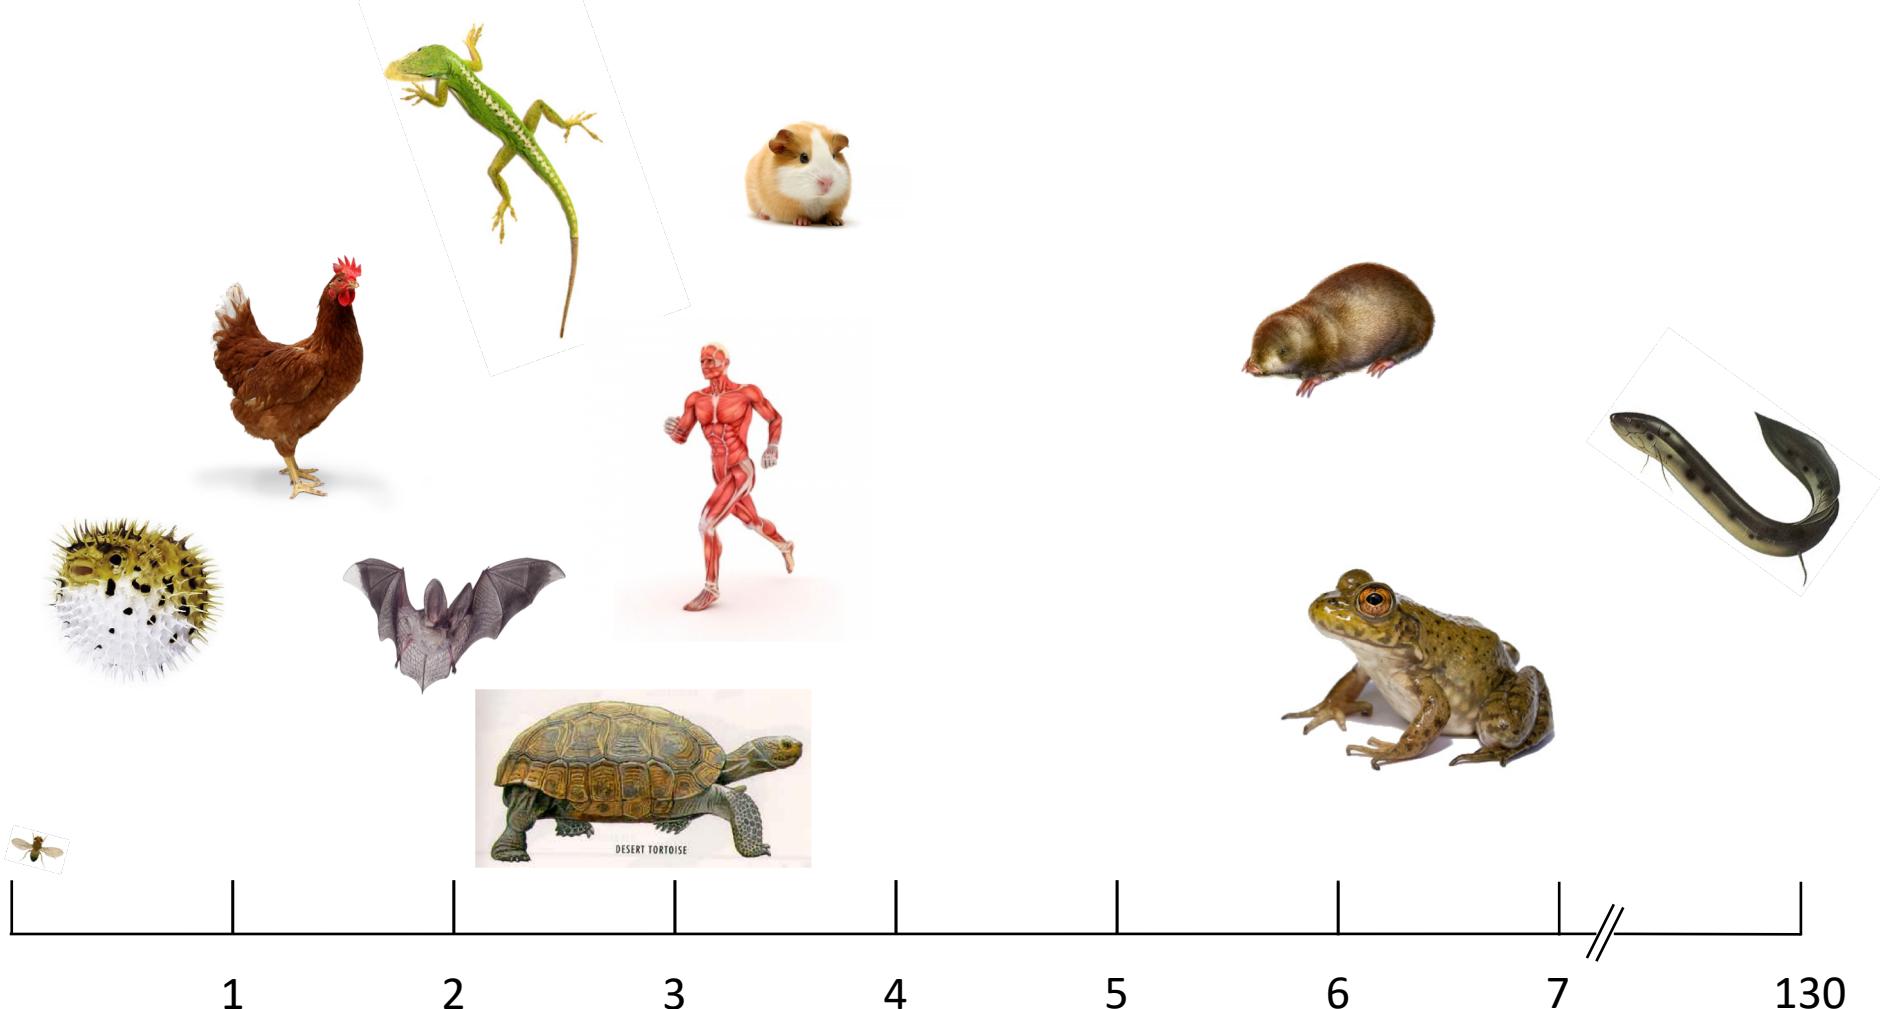


So You Want to Start a *de novo* Genome Assembly Project

Assuming you have a good reason to sequence and assemble a genome.

1. What is the size of the genome?
 2. What will be your sequencing “recipe”?
-
1. Do you have the computational resources?
 - *i.e.* a machine with 32 processors, 512GB RAM
 2. Do you have the time? Personnel? Bioinformatics experience?

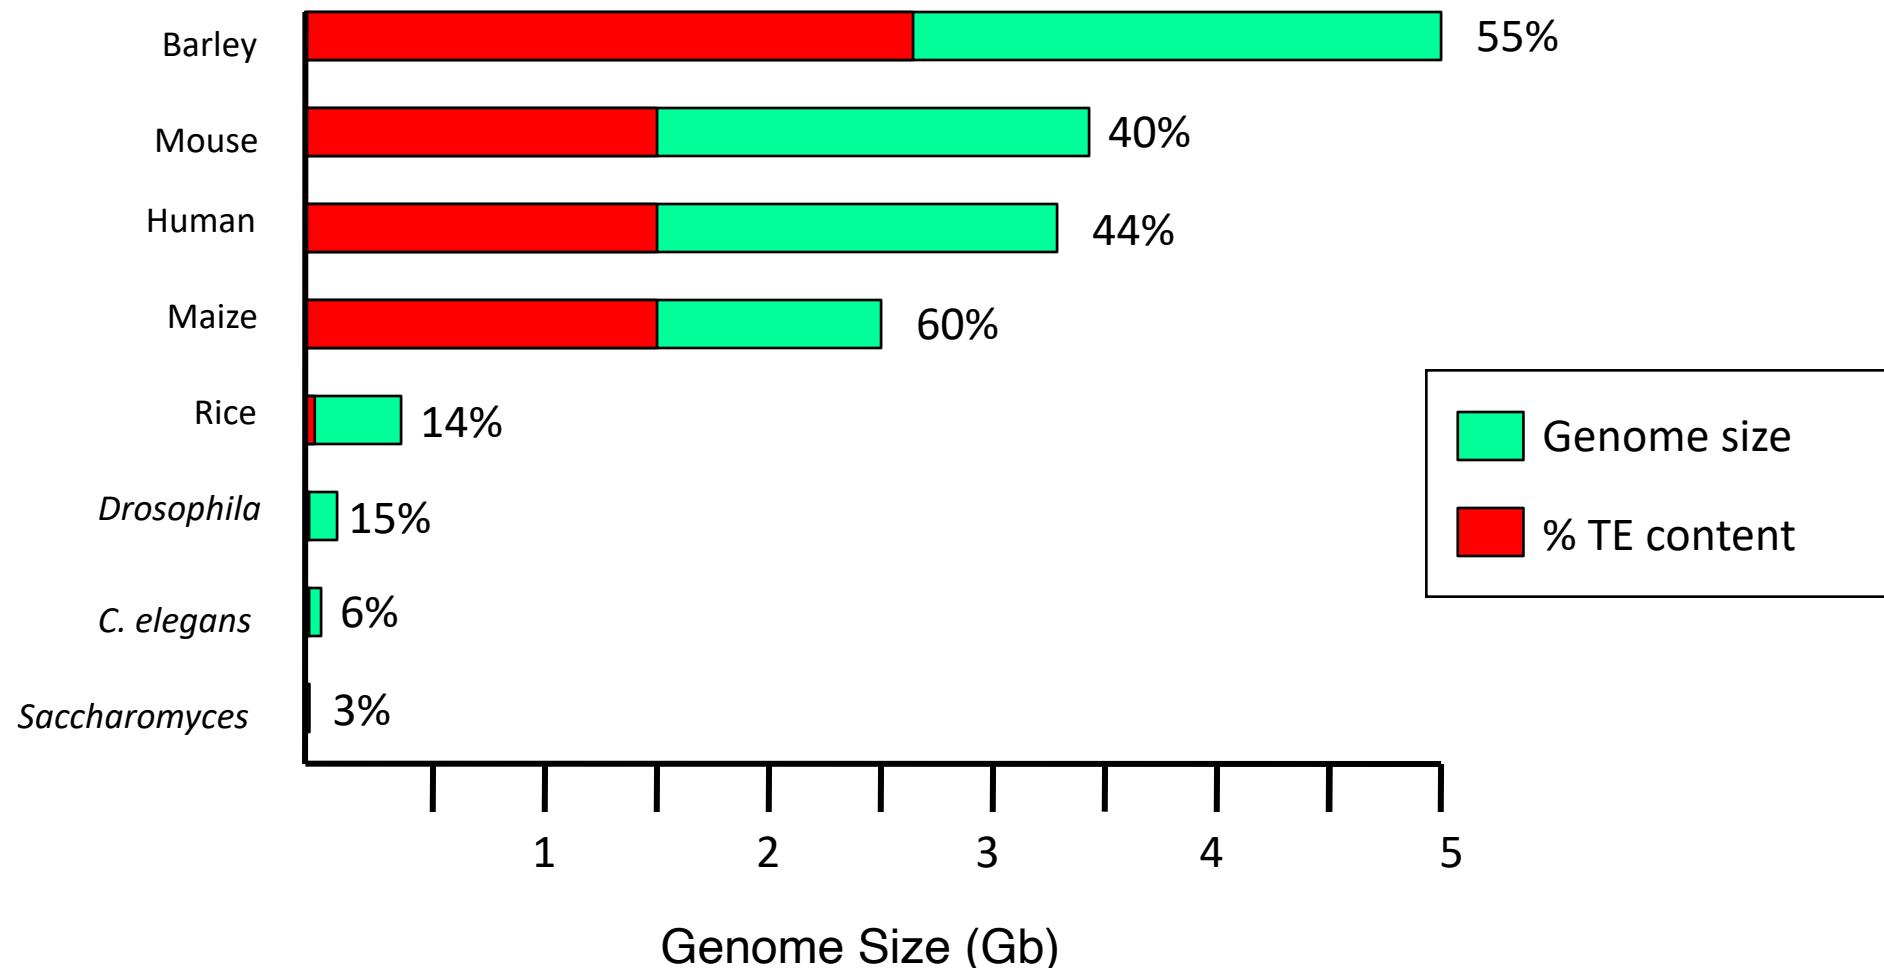
Animal Genome Sizes



Genome Size in Gigabases

<http://genomesize.com/>

Large genomes have a lot of transposable elements



Problem of Repeats

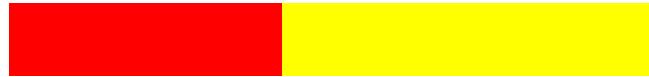
True sequence with repeats



Read 1



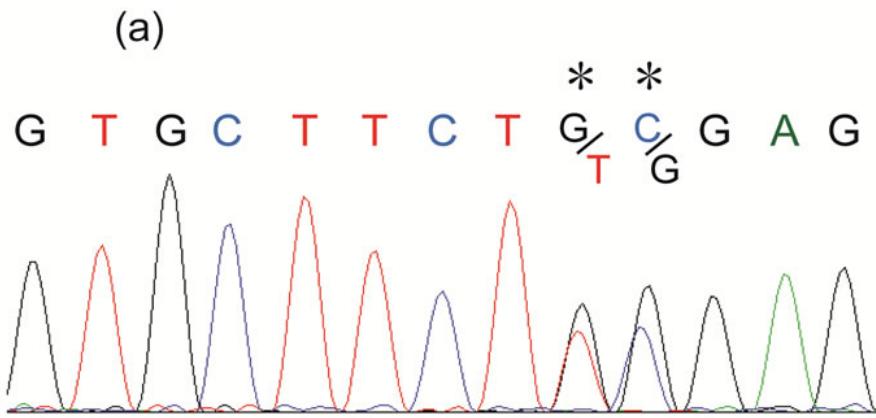
Read 2



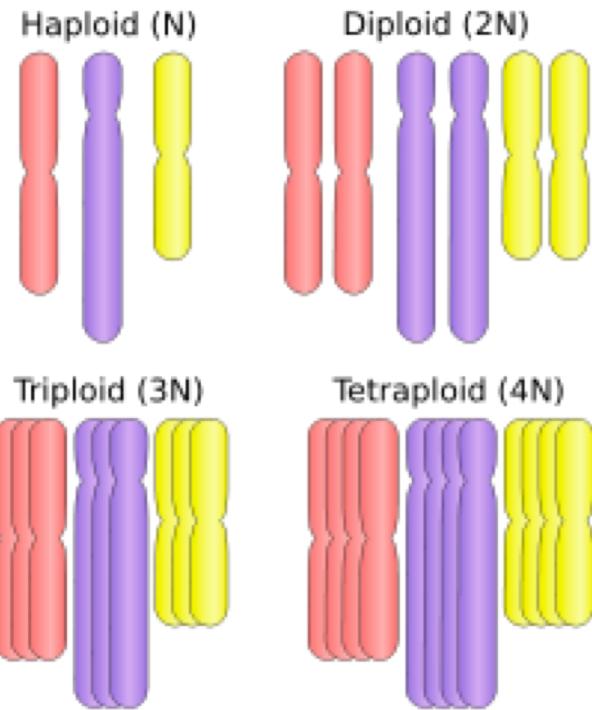
Overlap-consensus error has
omitted the green region



Other Questions



Expected Heterozygosity



Ploidy

Sequencing Technologies

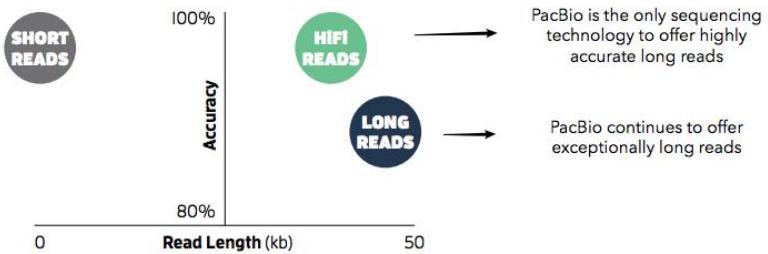
- Sanger method – old workhorse
 - “First generation sequencing”
 - Lower coverage, longer reads, fewer errors
- Next-generation sequencing – now standard
 - “Second generation”, PCR, short reads, more errors
 - 454, Illumina, SOLiD
- Third generation – becoming standard (expensive)
 - Single molecule, long reads, many more errors
 - PacBio, Oxford Nanopore

Basic Overview of Sequencing Technologies

	Company	Platform	Read Length	Time per run	Number of reads per unit	Common error
Short reads	Illumina	NovaSeq 6000	2 X 250bp	≤44 hours	≤40 billion	Substitution
	454 Life Sciences*	GS FLX Titanium XL	≤1000bp	700Mb/day	~150,000	indel
	Applied Biosystems*	SOLiD 5500xl W	25-50bp	5-8 days	~300 million	Substitution
Long reads	Pacific Biosciences	PacBio Sequel	≤60 kb (average 25 kb)	30 hours	≤4 million	indel
	Oxford Nanopore	MinION	≤4 Mb (average 6 kb)	≤72 hours	Tens of millions	indel

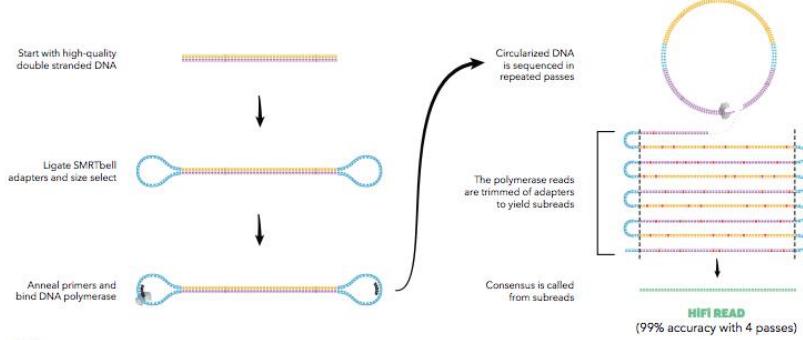
*retired

PacBio Sequel System



Generate Highly Accurate Long Reads

Produce HiFi reads using the circular consensus sequencing (CCS) mode to provide base-level resolution for detection of all variant types from single nucleotide to structural variants.



Optimize Your Run for Even Longer Reads

Sequence read lengths in the tens of kilobases using the continuous long read (CLR) sequencing mode to enable high-quality assembly of even the most complex genomes.

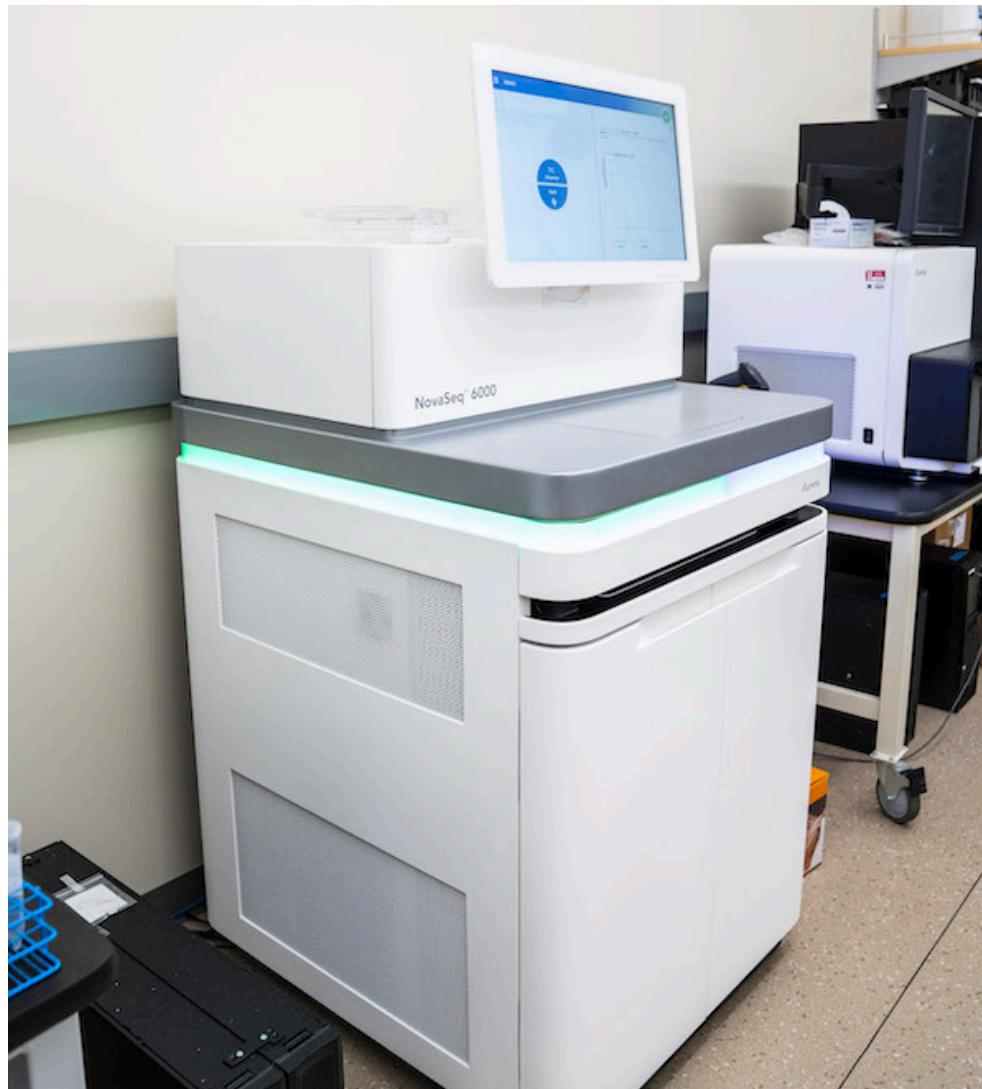
Half of Data in Reads
>50 kb

Longest Reads Up To
175 kb

Oxford Nanopore



illumina®



Illumina Paired-end and Mate-pairs

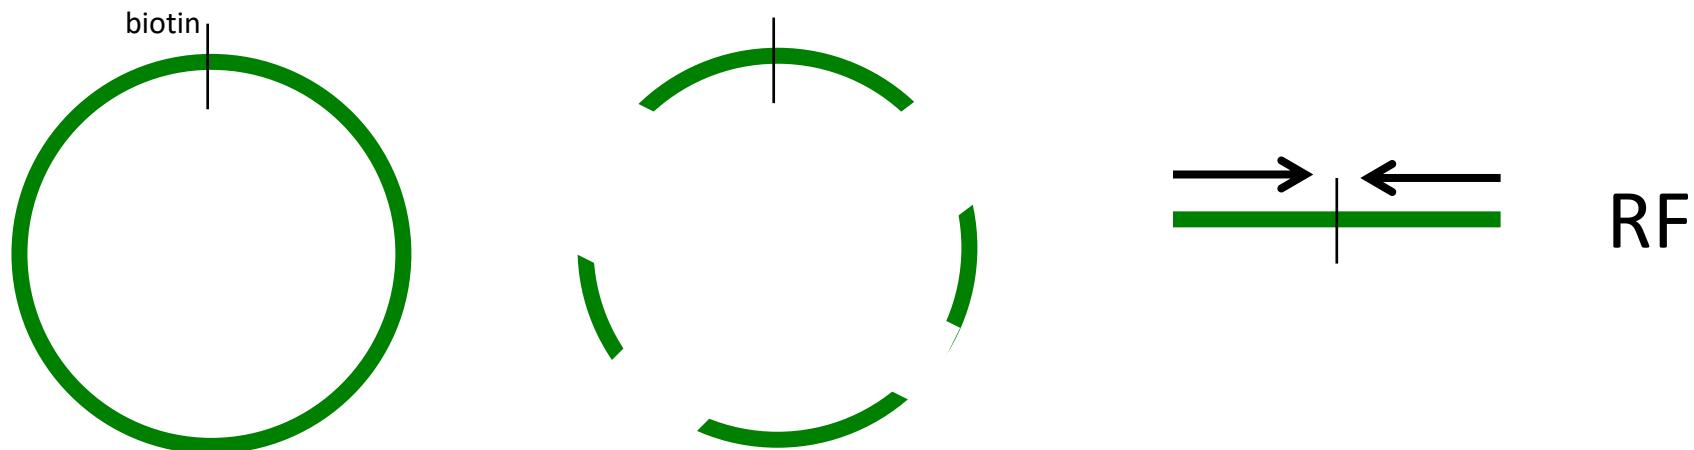
Paired-end (PE) “short insert library” sequencing

- Genome is fragmented to desired lengths
- Reads one end of the molecule, flips and then reads the other end
- Generates read pairs with a known distance between them



Mate-pair (MP) “jumping library” sequencing

- Circularizes longer molecules (2kb-25kb)
- Biotinylated, fragmented, enriched, and sequenced



Repeats Resolved

- Repeats can be resolved using paired-end information
 - If one end of a read is unique, then you can map both reads.
 - However, for longer repeats (*i.e.* LINEs) this will not work.
Unique Repetitive
 - Hence Illumina-based genomes tend to be fragmented.
- 

How much to sequence?

$$\text{Coverage} = \frac{(\text{read number} * \text{read length in bp})}{\text{Genome size in bp}}$$

(The number of times a site in the genome is represented by a read)

How much to sequence?

Say, a typical mammal...

$$100X = \frac{(2,000,000,000 * 150)}{3,000,000,000}$$

Libraries	Total Data (Gb)	Sequence Coverage
200bp paired-end	149.1	51.2X
500bp paired-end	141.7	48.7X
3kb mate-paired	57.3	19.7X
5kb mate-paired	72.5	24.9X
10kb mate-paired	28.5	9.8X
	449.1	154.3X

Cell Reports

Insights into the Evolution of Longevity from the Bowhead Whale Genome

Keane et al. (2015), *Cell Reports*

Resource



Fastq format



The screenshot shows a Mozilla Firefox browser window displaying a sample Fastq file. The file contains four entries, each representing a sequence read:

- @read1**
AGCTTATCCTCTGCTACCCCCGGGTTAGCGCACTTGATGTATTACAGC
+
BA1@CC7CBCCC9C8;B2@>C?B@B@B3=9?@B1:AB7B?B8B?B6B.7.
- @read2**
TTGGGCAGGATCTCCAGAACATATGGATGTGATCCACACAGCATTCTGC
+
?>?B@)<?@,AA7A@C<C?=@@B;+)?B5*@2=@+=BB,=B6C>AB@B24
- @read3**
TATGCTCAAGAAGGGGCTGATGAGTTGGTGTACGATATCACTGCCTC
+
A3AB:B1:B;9/0BBCBB<BB@AA0?BB9:BB<A@BB@7@6@<A@@@<3

- Each sequence has four lines
 1. sequence name starting with “@”
 2. Nucleotide sequence
 3. Empty line except for “+”
 4. Quality score information.

Phred Quality Scores

Logarithmically linked to error probabilities

Estimates several parameters based on peak shape and resolution at each base

These estimated parameters are then compared to lookup tables from known sequences

Allows the automation of quality control – especially helpful for large numbers of reads

Was invented for the Human Genome Project

Base-calling of automated sequencer traces using phred. II. Error probabilities. Ewing B, Green P (1998). Genome Research.

Base-calling of automated sequencer traces using phred. I. Accuracy assessment. Ewing B, Hillier L, Wendl MC, Green P. (1998). Genome Research.

Phred Quality Score	Probability of Error	Base Call Accuracy
10	1 in 10	90%
20	1 in 100	99%
30	1 in 1,000	99.9%
40	1 in 10,000	99.99%
50	1 in 100,000	99.999%
60	1 in 1,000,000	99.9999%

Data Quality Control

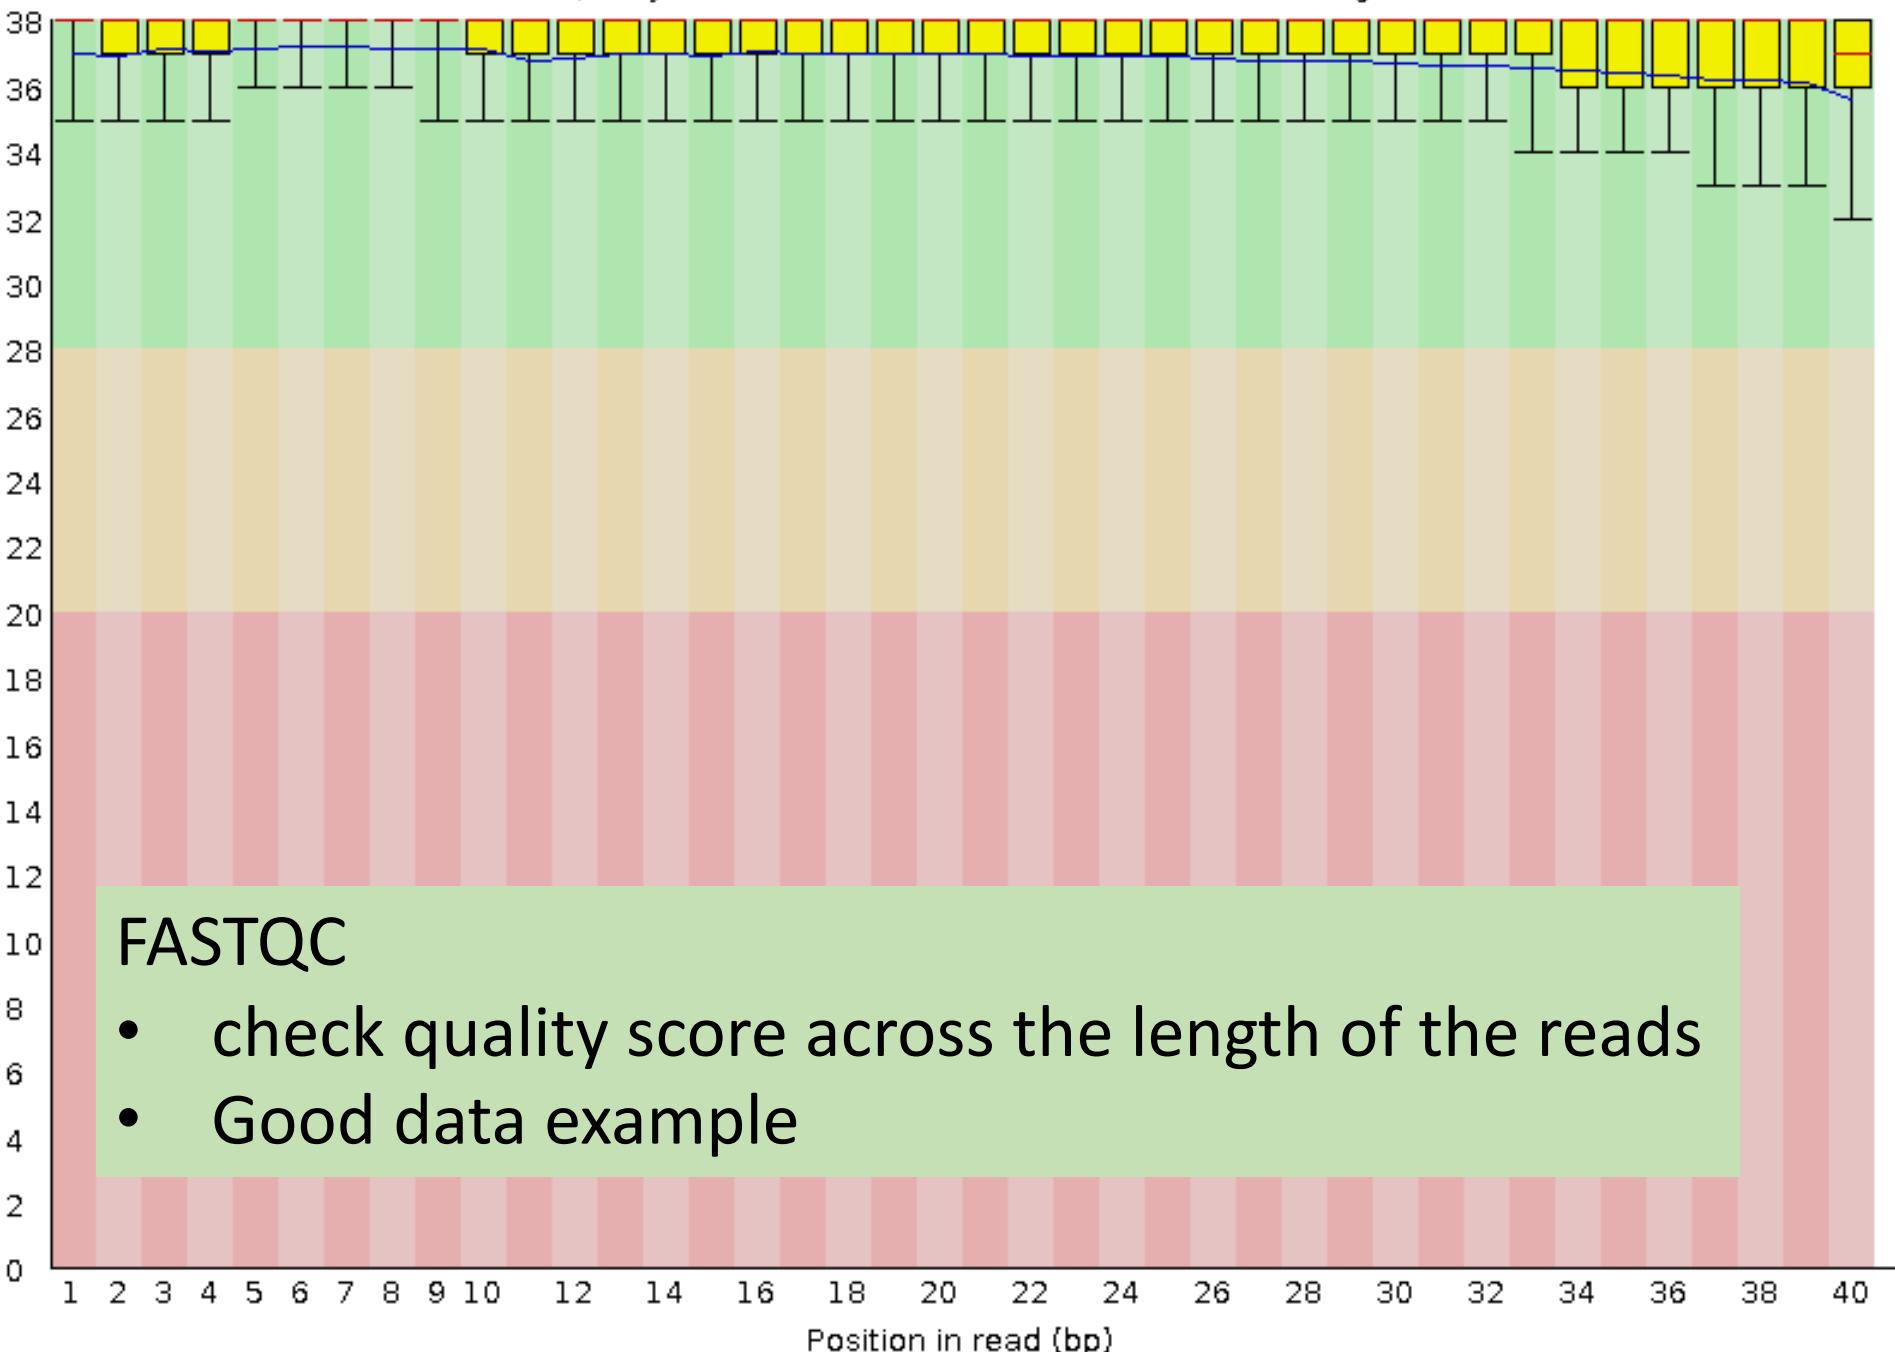
Quality score?

GC content?

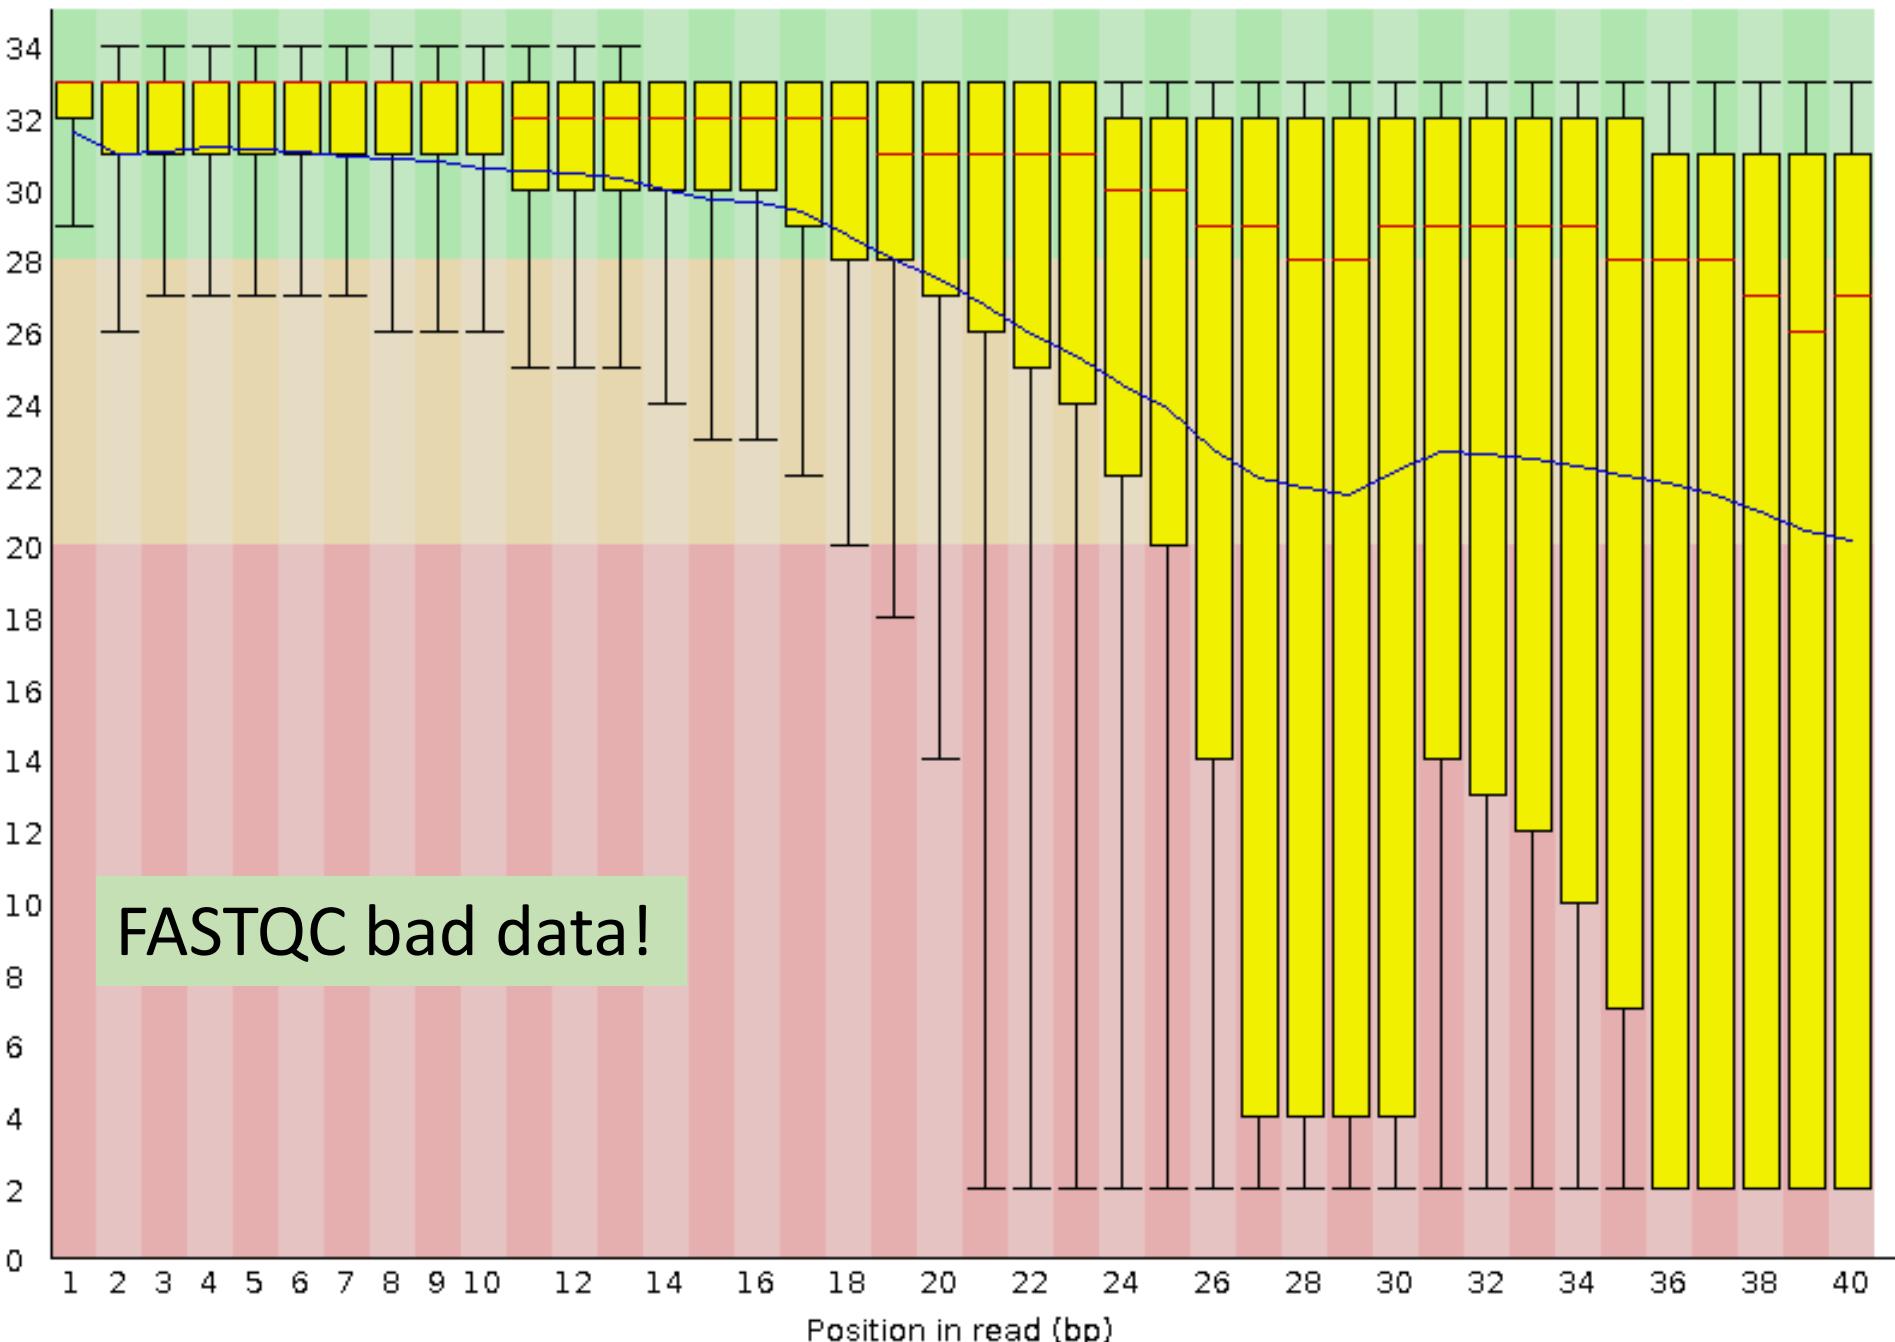
Sequence duplication levels?

Overrepresented sequences, or adaptors?

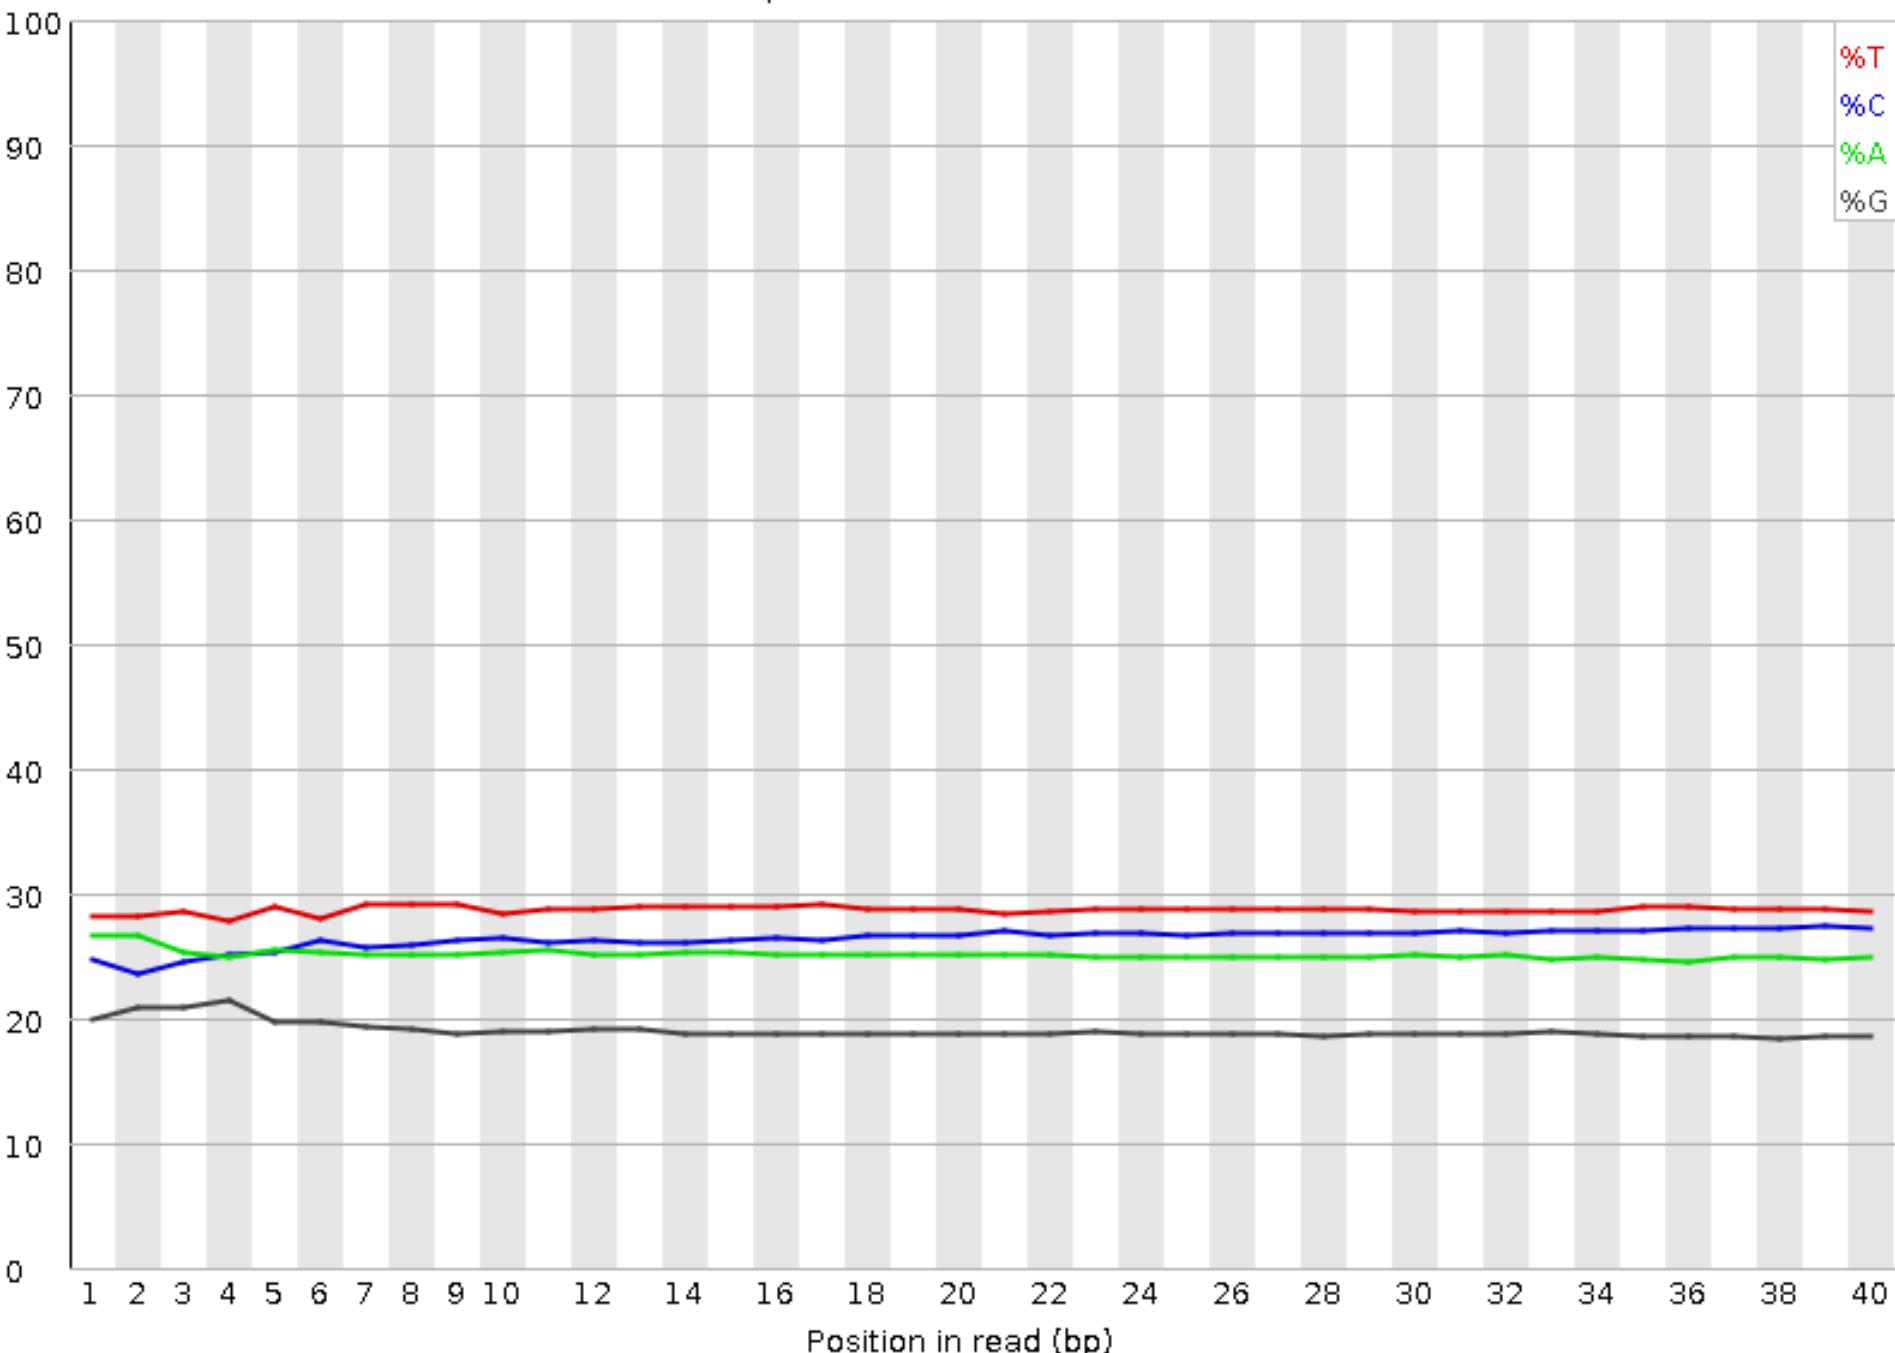
Quality scores across all bases (Illumina 1.5 encoding)



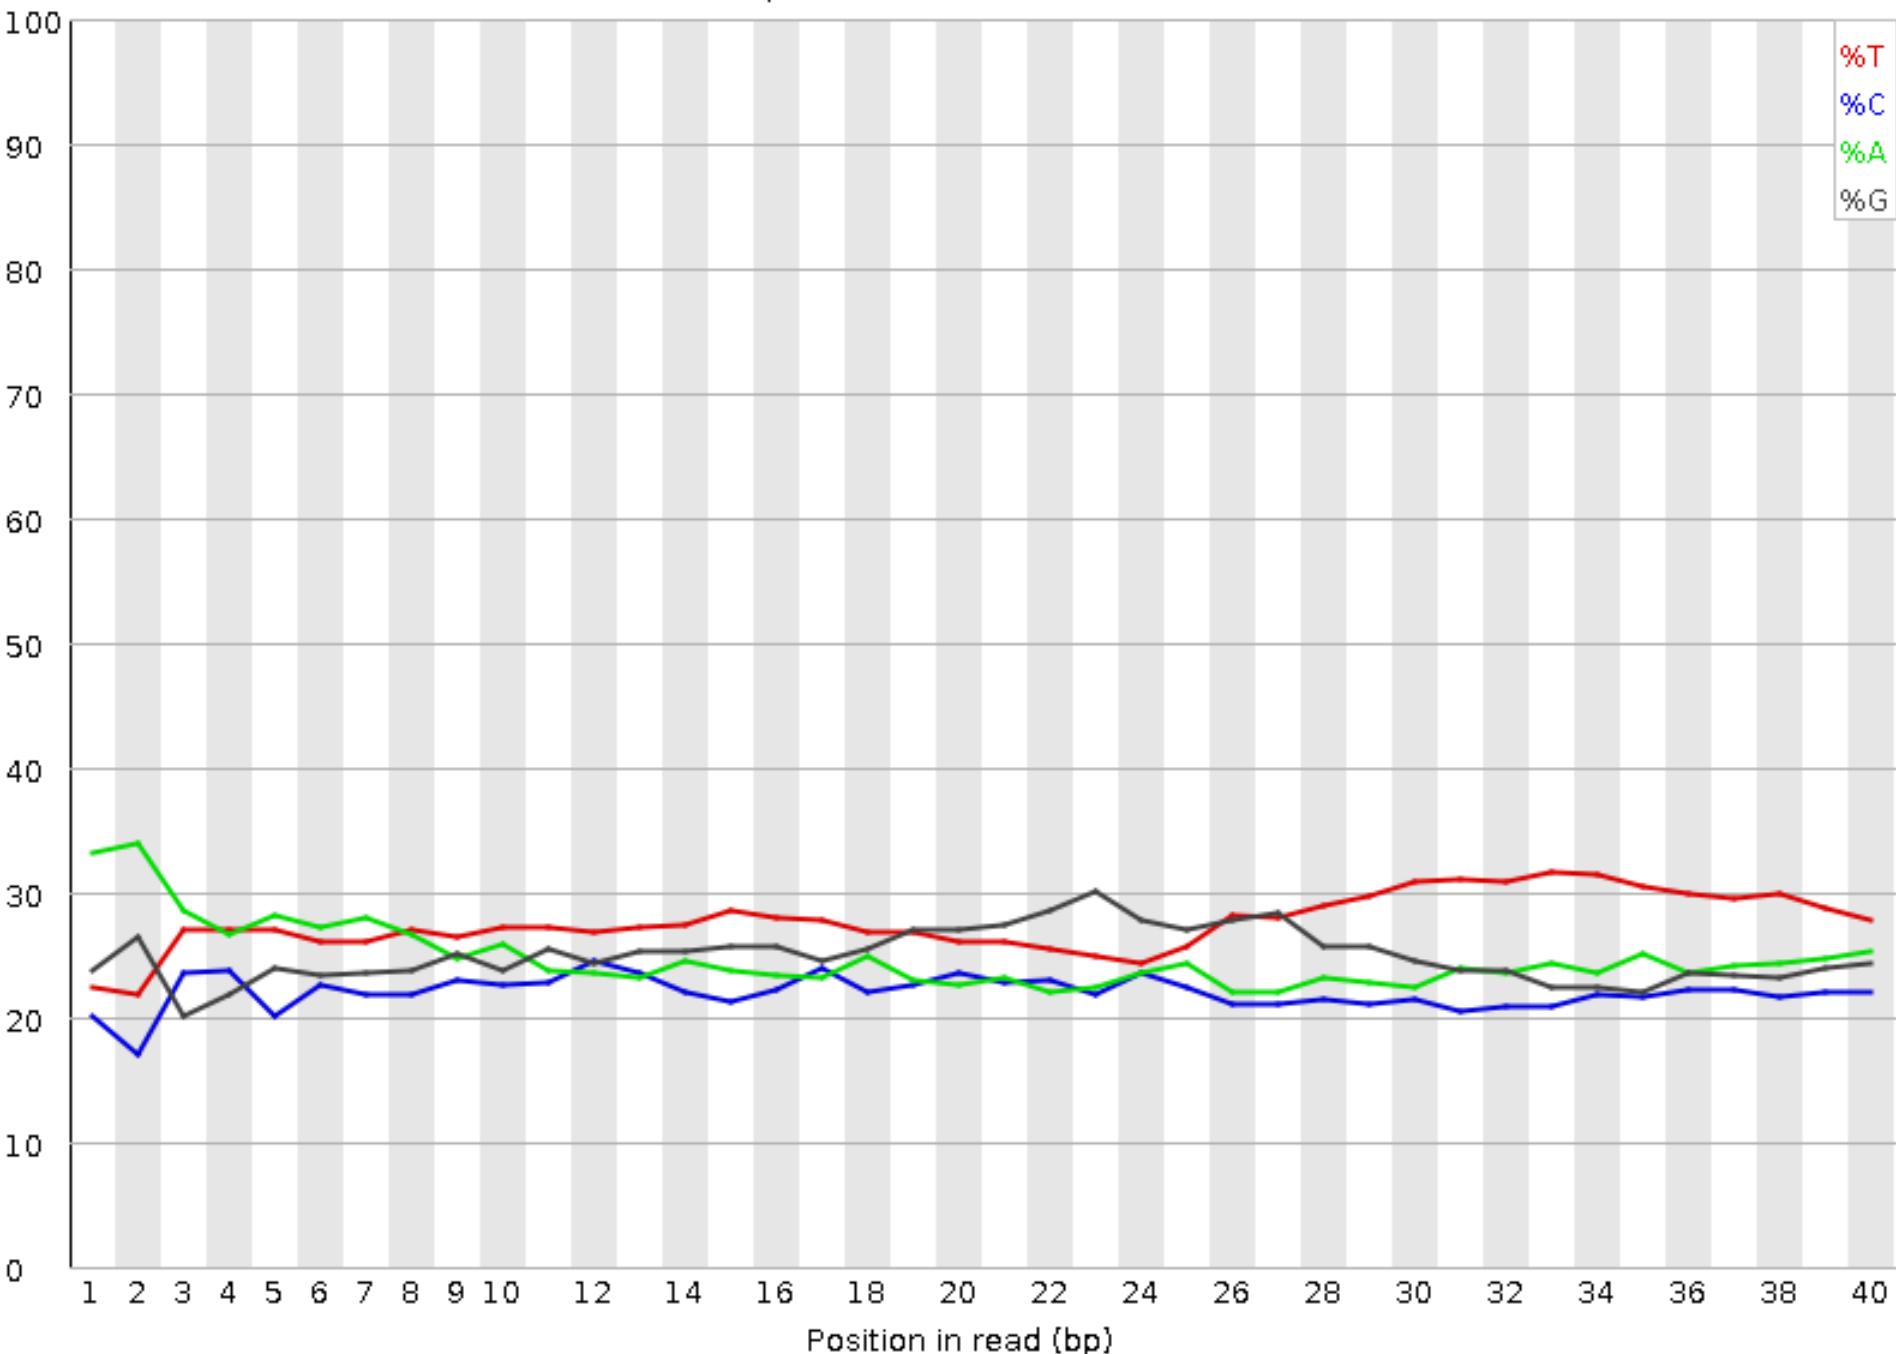
Quality scores across all bases (Illumina 1.5 encoding)



Sequence content across all bases



Sequence content across all bases



Trimming

- Trimmomatic (Bolger et al. 2014 *Bioinformatics*)
 - Paired end mode:
 - ILLUMINACLIP – trim adapters using database
 - HEADCROP
 - LEADING:3
 - TRAILING:3
 - SLIDINGWINDOW:4:15
 - MINLEN
 - Will trim entire dataset in pairs (F and R), and output singltons whose mates were eliminated

Three types of biases in NGS data

- Systematic bias
 - Problem with PCR, sequencer or library prep
 - Errors in base-calling
 - GC bias
 - High duplication rates
- Coverage bias
- Batch effects
 - Non-biological differences between experimental groups

Error correction

Substitution errors are most common in Illumina datasets

In theory, errors should be infrequent (0.0001) and random

But if you have 1 billion reads, that's 100,000 errors

Can be addressed by laying out all the reads covering a position

Use majority of reads to find (rare) erroneous sites, and correct them

Error correction

Error correction needs to be done before assembly!

K-spectrum-based correcting:

- Reads are broken up into k-mers

- Distances between k-mers are calculated

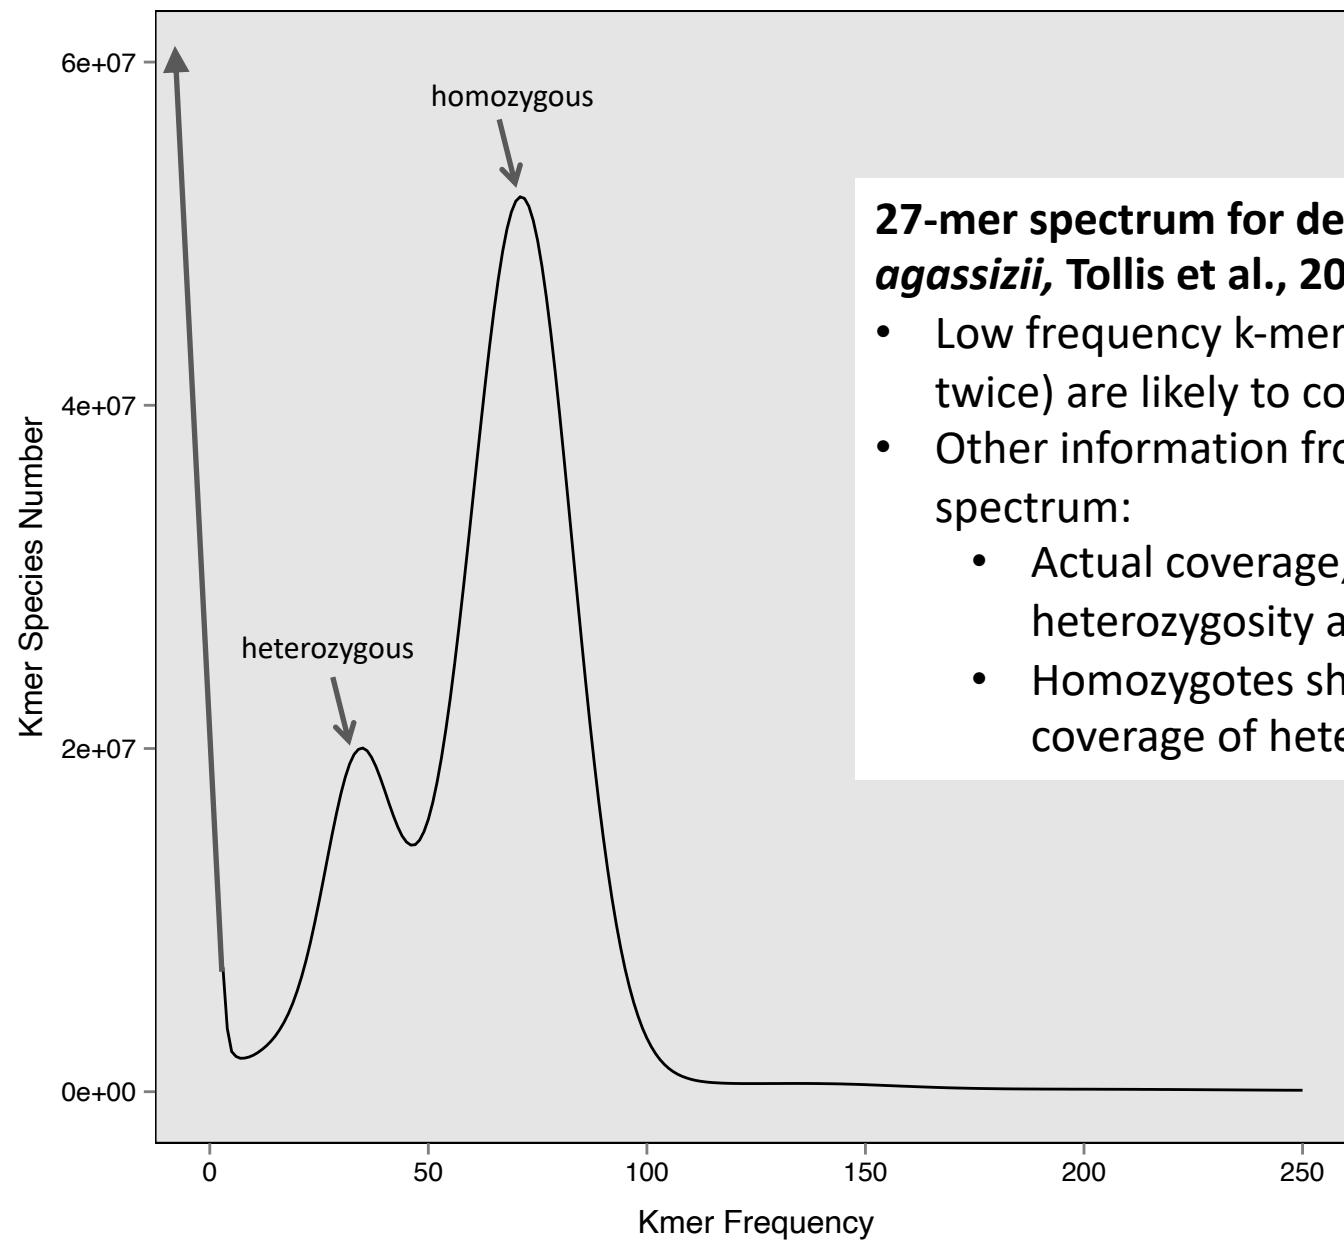
- Errors are corrected

Some tools are:

- SOAPdenovo error correction (Li et al. 2010)

- Quake (Kelley et al. 2010)

- Reptile (Yang et al. 2010)



27-mer spectrum for desert tortoise (*Gopherus agassizii*, Tollis et al., 2017 *PLoS One*)

- Low frequency k-mers (those occurring once or twice) are likely to contain errors
- Other information from k-mer frequency spectrum:
 - Actual coverage, given repetitiveness and heterozygosity across genome
 - Homozygotes should in theory be 2X coverage of heterozygotes

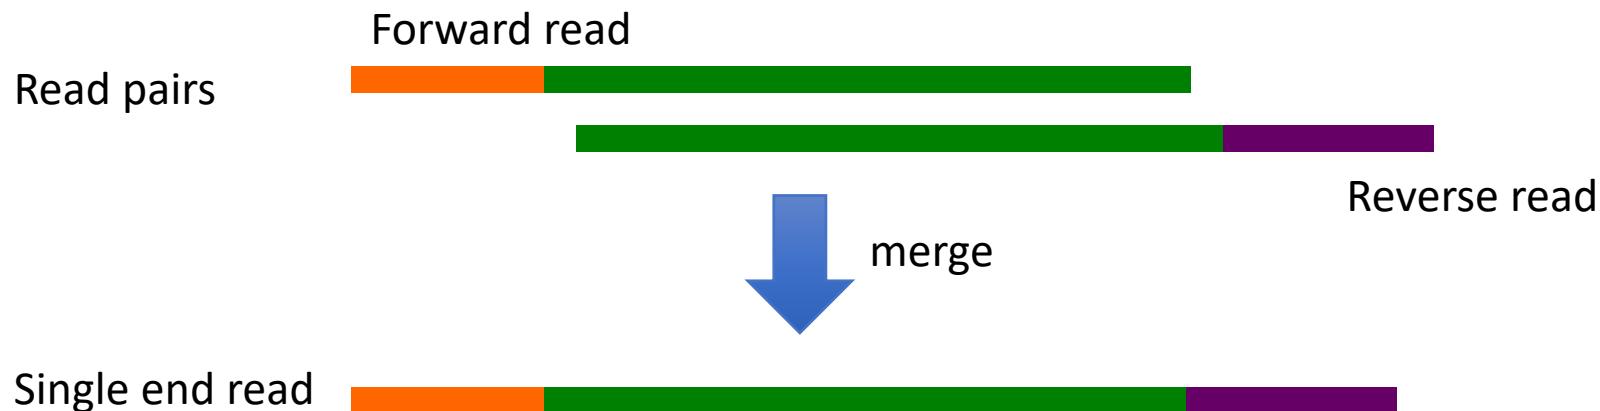
Merging reads

Joins together overlapping reads to create single-end reads

If you have 180/200bp libraries and ~100bp reads, this is recommended.

Speeds up assembly considerably so the assembler does not have to calculate the distance between these reads.

FLASH, Cope, PEAR are some tools that do this.



Choose genome,
gather info



DNA Library
preparation



Sequencing



Quality check



Trimming



Error correction



Merge
overlapping reads



ASSEMBLE!

